

Chawin Sitawarin

Soda Hall, UC Berkeley
Berkeley, CA 94709
☎ 510-423-2880
✉ chawins@berkeley.edu
🌐 <https://chawins.github.io/>

*PhD Candidate, Computer Science @ UC Berkeley
Machine Learning Security & Privacy Researcher*

Education

- 2018–2024 **PhD in Computer Science**, *UC Berkeley*, Berkeley CA.
Advisor: Professor David Wagner | GPA 3.86
- 2014–2018 **BSE in Electrical Engineering (High Honor)**, *Princeton University*, Princeton NJ.
Cumulative GPA: 3.90, Departmental GPA: 3.95 | Certificate in Applications of Computing

Research Interests

I am broadly interested in the security and safety aspects of ML systems. My current focus is on security and privacy issues of large language models. My past works are primarily on the adversarial robustness of vision systems. My research goal is to make ML models secure, robust, and trustworthy in practice without compromising their utility.

Selected Publications

- 2023 **Defending against Transfer Attacks from Public Models**, [C. Sitawarin](#), J. Chang*, D. Huang*, W. Altoyan, D. Wagner, Preprint (Under Submission), [Paper](#), [Code](#).
- 2023 **OODRobustBench: Benchmarking and Analyzing Adversarial Robustness Under Distribution Shift**, L. Li, Y. Wang, [C. Sitawarin](#), M. Spratling, Preprint (Under Submission), [Paper](#).
- 2023 **REAP: A Large-Scale Realistic Adversarial Patch Benchmark**, N. Hingun*, [C. Sitawarin](#)*, J. Li, D. Wagner, Preprint (Under Submission), [Paper](#), [Code](#).
- 2023 **Preprocessors Matter! Realistic Decision-Based Attacks on Machine Learning Systems**, [C. Sitawarin](#), F. Tramèr, N. Carlini, Preprint (Under Submission), [Paper](#), [Code](#).
- 2023 **Part-Based Models Improve Adversarial Robustness**, [C. Sitawarin](#), K. Pongmala, Y. Chen, N. Carlini, D. Wagner, ICLR 2023 (Poster), [Paper](#), [Code](#).
- 2022 **Demystifying the Adversarial Robustness of Random Transformation Defenses**, [C. Sitawarin](#), Z. Golan-Strieb, D. Wagner, ICML 2022 (Short Presentation) and AAAI-22 AdvML Workshop (Best Paper), [Paper](#), [Code](#).
- 2021 **Adversarial Examples for k -Nearest Neighbor Classifiers Based on Higher-Order Voronoi Diagrams**, [C. Sitawarin](#), E. M. Kornaropoulos, D. Song, D. Wagner, NeurIPS 2021 (Poster), [Paper](#), [Code](#).
- 2021 **Improving the Accuracy-Robustness Trade-Off for Dual-Domain Adversarial Training**, [C. Sitawarin](#), A. Sridhar, D. Wagner, Workshop on Uncertainty & Robustness in Deep Learning (ICML 2021), [Paper](#), [Code](#).
- 2021 **Mitigating Adversarial Training Instability with Batch Normalization**, A. Sridhar, [C. Sitawarin](#), D. Wagner, Workshop on Security and Safety in Machine Learning Systems (ICLR 2021), [Paper](#).
- 2021 **SAT: Improving Adversarial Training via Curriculum-Based Loss Smoothing**, [C. Sitawarin](#), S. Chakraborty, D. Wagner, AISC 2021 (co-located with CCS), [Paper](#).
- 2020 **Minimum-Norm Adversarial Examples on k -NN and k -NN-Based Models**, [C. Sitawarin](#), D. Wagner, Deep Learning and Security Workshop (IEEE S&P 2020), [Paper](#).
- 2019 **Analyzing the Robustness of Open-World Machine Learning**, V. Sehwag, A. N. Bhagoji, L. Song, [C. Sitawarin](#), D. Cullina, M. Chiang, and P. Mittal, AISC 2019 (co-located with CCS), [Paper](#).
- 2019 **Defending Against Adversarial Examples with K-Nearest Neighbor**, [C. Sitawarin](#), D. Wagner, Preprint, [Paper](#), [Code](#).

- 2018 **On the Robustness of Deep k-Nearest Neighbors**, [C. Sitawarin](#), D. Wagner, Deep Learning and Security Workshop (IEEE S&P 2019), [Paper](#).
- 2018 **Not All Pixels are Born Equal: An Analysis of Evasion Attacks under Locality Constraints**, V. Sehwag, [C. Sitawarin](#), A. N. Bhagoji, A. Mosenia, M. Chiang, P. Mittal, CCS 2018 Poster, [Paper](#).
- 2018 **DARTS: Deceiving Autonomous Cars with Toxic Signs**, [C. Sitawarin](#), A. N. Bhagoji, A. Mosenia, M. Chiang, P. Mittal, Preprint, [Paper](#).
- 2018 **Rogue signs: Deceiving Traffic Sign Recognition with Malicious Ads and Logos**, [C. Sitawarin](#), A. N. Bhagoji, A. Mosenia, M. Chiang, P. Mittal, Deep Learning and Security Workshop (IEEE S&P 2018), [Paper](#).
- 2018 **Enhancing Robustness of Machine Learning System via Data Transformations**, A. N. Bhagoji, D. Cullina, [C. Sitawarin](#), P. Mittal, CISS 2018, [Paper](#).

Other Experiences

- Summer 2022 **Google, Sunnyvale**, Research Intern.
Evaluate and mitigate machine learning security risks in a practical setting where a pair of public client-side and secret server-side models is deployed for a malware detection task. Hosted by Ali Zand and David Tao.
- Fall 2021 - **Google, Remote**, Student researcher (part-time).
- Spring 2022 Developed threat model and appropriate evaluation for adversarial robustness in new and practical settings (e.g., dynamic models, black-box model recovery). Hosted by Nicholas Carlini.
- Summer 2021 **Nokia Bell Labs, Remote**, Summer research intern.
Investigated relationships between causality and robustness in machine learning, focusing on leveraging causal relationship to improve robustness and generalization to unseen attacks/corruptions. Mentored by Anwar Walid.
- Fall 2020 **EECS Department, UC Berkeley, Berkeley CA**, Graduate student instructor.
Part of the content development team for CS189/289A: Introduction to Machine Learning. Created homework problems and materials for the discussion sections and taught discussion sections.
- Summer 2019 **IBM Research, Yorktown Heights NY**, Summer research intern.
Studied the effectiveness of existing defenses against adversarial examples from a perspective of concentration bound and improved adversarial training through optimization techniques. Mentored by Supriyo Chakraborty.

Awards & Honors

- | | | |
|-----------|--|--|
| 2022 | Google-BAIR Commons Project | <i>Research grant</i> |
| 2021-2022 | Center for Long-Term Cybersecurity (CLTC) | <i>Research grant</i> |
| 2021 | Microsoft-BAIR Commons Project | <i>Research grant</i> |
| 2018 | Phi Beta Kappa | <i>Academic Honor Society</i> |
| 2018 | Sigma Xi | <i>Scientific Research Honor Society</i> |
| 2017 | The P. Michael Lion III Fund | <i>Summer research funding for Princeton engineering students</i> |
| 2016 | Tau Beta Pi | <i>Engineering Honor Society</i> |
| 2016 | Shapiro Prize for Academic Excellence | <i>Academic award at Princeton University</i> |
| 2013 | King's Scholarship | <i>Prestigious scholarship awarded by Thai government for pursuing a bachelor's degree</i> |

Activities and Services

- Program Committee, AISEC '22, '23.**
- Reviewer, ICLR '24 | ICML '22 (top reviewer) | NeurIPS '22, '23 | BAND workshop '23.**
- 2019-present **DARE: Diversifying Access to Research in Engineering**, *Mentor*, I have mentored multiple students from DARE, a program that promotes diversity in EECS undergraduate research.
- 2018-2020 **CSGSA, Treasurer**, Computer Science Graduate Student Assembly at UC Berkeley.
- 2018-2019 **Security Seminar, Organizer**, Organized a biweekly seminar on security and privacy at UC Berkeley, hosting outside speakers from both industry and academia.